
From a traditional to an AI-native SDLC

AI is new, and everyone wants it.

The tools change every month, and the pressure to adopt them comes from every direction at once. What is missing is a direction of our own.

Everyone steers, nobody drives

Every team picks its own tools, its own prompts and its own workflow. The result is fragmentation, and the direction ends up set by whoever sells the tool.

Illustration: Jan Meskens

You are not steering. You are being driven.



What we bring

Direction

One way of working that every team shares, instead of a tool and a habit per team.

Expertise

People who work with agents every day and know which practices hold up on real code.

Technical know-how

Hands on in your repositories, pipelines and tooling, not only on slides.

Tom Haeldermans



Founder · SmartAgents

PhD

In industrial engineering (and dropped it for IT)

On holiday

Found somewhere on a mountain

Marathon

Ran one, thinking about the next

Bas Sarrechia

Founder • SmartAgents

Chemistry


A goat cheese factory, the SCK CEN nuclear research centre, plastics, and more

Obsessive

Easily, about whatever has his attention

Currently

Fitness, AI and the stock market



Issues we saw in our first weeks

Belief

Some distrust the output entirely, others expect miracles. It is the same mistake in opposite directions.

Capability

No training, too little access to AI tooling, and processes built for a world where writing code was the expensive part.

Foundation

Product knowledge lives in people's heads. AI workflows are fragmented, each only as good as whoever built it, and everything is solved with a skill.

Adoption

The spread between teams is huge: some work almost AI-first, others barely use AI. There is no common way of working to close that gap.

The hardest constraint is not technical

"We've always done it this way"

Answered by someone inside the team who shows, on real work, that it can be done differently.

"I don't trust it. My code is better"

Answered by real changes that show the output meets your standards, once the context and the gates are set up well.


"It only creates slop"

Answered by gates. Slop is what unchecked output looks like. Most teams have review and CI, but not every team applies them consistently.

The bottleneck moves. It does not disappear.


Today

Six desks, five handoffs



AI-assisted

What you ask and who checks it



AI-native

One run, four decisions



Four gates 1 Intent accepted 2 Spec approved 3 Change approved 4 Release authorised

Each row narrows where work waits: Build today, then Analysis and Review, then only the release sign-off.

At each gate a person decides, and the pipeline makes sure no gate is skipped.

Adapted from Anthropic's AI-Native SDLC playbook.

Realisations after 1 month at Ponto

The first quality gate: a reusable test setup for the Financial Institute Connectors (FICs).

FICs covered by tests

16 → 99 all of them

Code coverage

19% → 100%

Along the way

Readable tests

The team is happy with how the tests read.

FIC configurations fixed

2 bad configurations removed, 2 others updated.

17 production issues

Identified, not yet resolved.

Dead code paths

Several identified, not yet resolved.

What the quality gates bring

No rogue AI

A gate limits the blast radius of a change, so an agent cannot quietly modify more than it should.

More trust in AI

What an agent may change is limited and checked, so its output can be trusted further.

Better output

Constraints push an agent towards higher-quality work instead of whatever passes first.

Standardisation

Gates enforce standards such as code coverage, and the same mechanism works for security.

Next objectives (4 weeks)

Follow-ups

Clean up the findings

Resolve the production issues and dead code paths found this month.

Logging strategy

Direct impact: less load on Kibana, and it helps other issues surface along the way.

Team coaching

Based on the team’s own /insights from Claude.

The big one

A proposal for requirement 2



Catch

Issues arrive by email as they happen.

Ingest


A website takes them in.

Specify

Requirements and acceptance criteria: gate 2.

Builds on the quality gate, and is needed for the implementation proposal for requirement 2.

Inside Build: the agent runs the gates itself



The gates hold every change to today's standards, or stricter, before a person ever sees it.

Thank you

Questions or concerns?

smartagents.be